

# Software and Data Engineering CSL 7090

*Assignment 3: Deploying Application to Server*  
*Instructor: Dr. Sumit Kalra*



॥ त्वं ज्ञानमयो विज्ञानमयोऽसि ॥

Submitted by

Ajit Kumar  
M21CS017

# Introduction

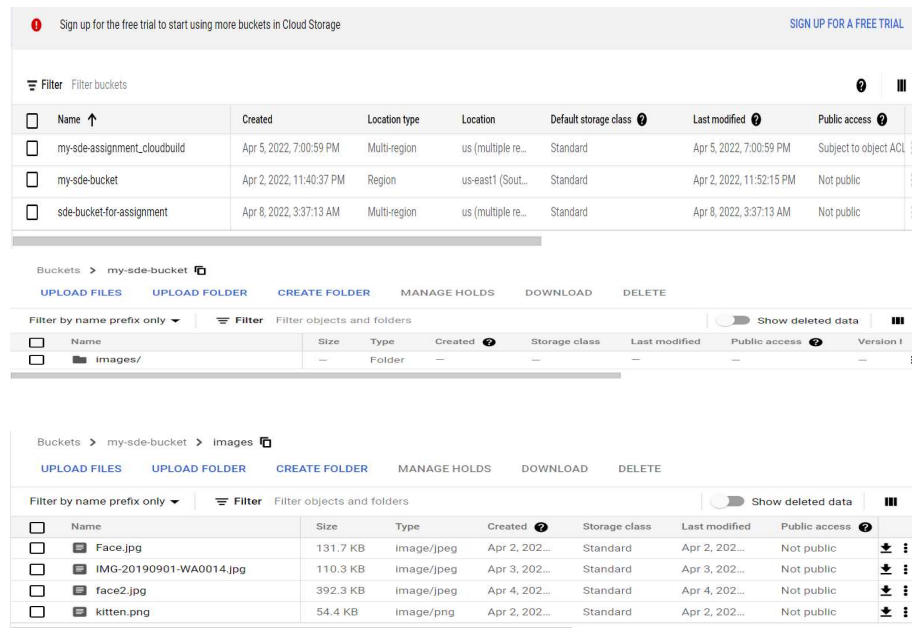
I have written a program in python for face detection using GOOGLE CLOUD VISION API. It takes the image from the current directory and google cloud storage too and detect the face inside the image and draw line around the faces.

To deploy the service , I have used kubernetes engine. In kubernetes engine deployment, I have make **dockerimage** . Docker image is stored into the Artifact registry. Finally I deploy this service through kubernetes to the google cloud.

For load balancing,

# Procedure

- First created a project as **my-sde-assignment**.
- Then I have created a bucket to upload all the image that is used in our program into cloud storage bucket name as **my-sde-bucket**.



- Then I have written program for my requirement.
- Below screen shot is program for downloading the image from cloud storage(bucket).

```
def download_blob(bucket_name, source_blob_name, destination_file_name):
    """Downloads a blob from the bucket."""
    # The ID of your GCS bucket
    # bucket_name = "your-bucket-name"
    # The ID of your GCS object
    # source_blob_name = "storage-object-name"
    # The path to which the file should be downloaded
    # destination_file_name = "local/path/to/file"
    storage_client = storage.Client()
    bucket = storage_client.bucket(bucket_name)
    # Construct a client side representation of a blob.
    # Note `Bucket.blob` differs from `Bucket.get_blob` as it doesn't retrieve
    # using `Bucket.blob` is preferred here.
    blob = bucket.blob(source_blob_name)
    blob.download_to_filename(destination_file_name)
    print(
```

- Below screen shot is program for detect face.

```
# [START vision_face_detection_tutorial_send_request]
def detect_face(face_file, max_results=10):
    """Uses the Vision API to detect faces in the given file.

    Args:
        face_file: A file-like object containing an image with faces.

    Returns:
        An array of Face objects with information about the picture.
    """
    # [START vision_face_detection_tutorial_client]
    client = vision.ImageAnnotatorClient()
    # [END vision_face_detection_tutorial_client]

    content = face_file.read()
    image = vision.Image(content=content)
```

- To select the vision API in google cloud:
  - a. First select the my project as my-sde-assignment.
  - b. Next clone the sample repository
  - c. To use the vision API, our app need to authenticate its identity to the vision service.

d. Created a service account to our API request.

The screenshot shows the Google Cloud IAM & Admin console. On the left, the 'IAM & Admin' menu is expanded, showing options like Billing, APIs & Services, Marketplace, Compute Engine, Cloud Storage, VPC network, Kubernetes Engine, BigQuery, and SQL. On the right, a list of IAM roles is displayed, including IAM, Identity & Organization, Policy Troubleshooter, Policy Analyzer, Organization Policies, Service Accounts, Workload Identity Federation, Labels, Tags, Settings, Privacy & Security, Identity-Aware Proxy, Roles, Audit Logs, Manage Resources, Create a Project, Asset Inventory, and Essential Contacts.

Below the menu, a section titled 'Grant this service account access to project (optional)' is shown. It contains a text box with the role 'Owner' selected from a dropdown menu. Below the role, it says 'Full access to most Google Cloud resources. See the list of included permissions.' There is a link to 'Add condition' and a button to 'ADD ANOTHER ROLE'. A 'CONTINUE' button is at the bottom of the section.

Below this section, a heading 'Create private key for "vision-detect-2"' is shown. It includes a warning: 'Downloads a file that contains the private key. Store the file securely because this key can't be recovered if lost.' There are two radio buttons for 'Key type': 'JSON' (selected) and 'P12'. Below 'JSON' is the text 'Recommended'. Below 'P12' is the text 'For backward compatibility with code using the P12 format'. At the bottom right, there are 'CANCEL' and 'CREATE' buttons.

e. Then I created service account key in same directory as vision-detect-2 and set the key as default credentials. To do this I have used following command:

```
gcloud iam service-accounts keys \
create key.json --iam-account \
vision-detect-2@my-sde-assignment.iam.gserviceaccount.com
```


```
export GOOGLE_APPLICATION_CREDENTIALS=key.json
```

- Then to install all the requirement for the program from requirement.txt ,used the following command:

**pip3 install -r requirements.txt**


- Then finally run the program:
  - a. This is taking the image from the same directory:

```
]
kumar_281@cloudshell:~/.../snippets/face_detection (my-sde-assignment)$ python3 faces.py ./re
Found 1 face
Writing to file out.jpg
kumar_281@cloudshell:~/.../snippets/face_detection (my-sde-assignment)$ █
```



- b. This is image taking from the google cloud storage:

```
kumar_281@cloudshell:~/.../snippets/face_detection (my-sde-assignment)$ python3 faces.py --max-results 20 gs:/
0190901-WA0014.jpg
Downloaded storage object images/IMG-20190901-WA0014.jpg from bucket my-sde-bucket to local file ./resources/IM
Found 9 faces
Writing to file out.jpg
kumar_281@cloudshell:~/.../snippets/face_detection (my-sde-assignment)$ █
```



## Deployment

- For the deployment I have used the kubernetes , So I have created docker file.

```

Dockerfile
1 FROM ubuntu:21.04
2
3 ENV TZ=Asia/Kolkata \
4     DEBIAN_FRONTEND=noninteractive
5
6
7 COPY resources /exp/resources
8 COPY faces.py /exp/faces.py
9 COPY key.json /exp/key.json
10 COPY requirements.txt /exp/requirements.txt
11 #RUN apt update && apt install -y tcl
12 RUN apt-get update && apt-get install -y python3.9 python3-pip
13 WORKDIR /exp
14 RUN pip3 install -r /exp/requirements.txt
15 RUN export GOOGLE_APPLICATION_CREDENTIALS=key.json
16 # RUN python3 faces.py /resources/faces-input img

```

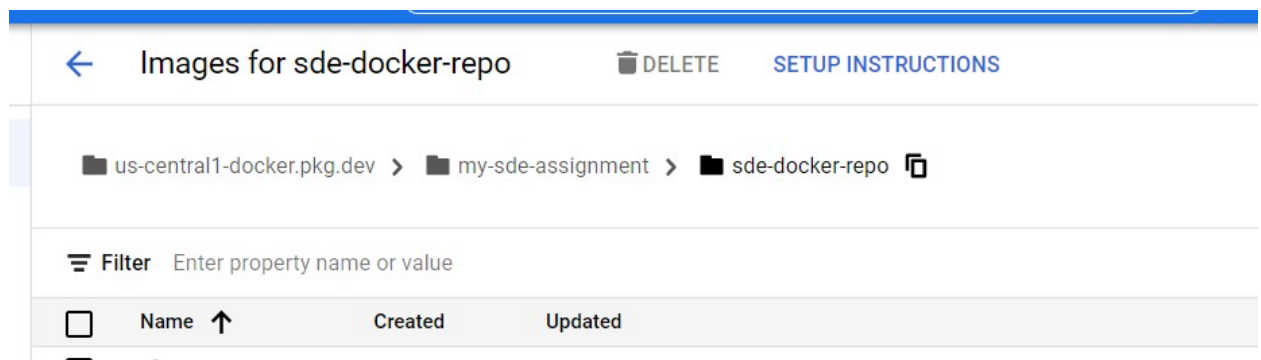
- Then I created image for above docker file:
  - a. First I have created a artifact registry name as sde-docker-repo by using following command

```

gcloud artifacts repositories create sde-docker-repo \
  --repository-format=docker \
  --location=us-central1 \
  --description="Docker repository"

```

- b. Next, I have created docker image by using following command:  
**docker build -t REGION-docker.pkg.dev/\${my-sde-assignment}/sde-docker-repo/sde-doc-image:latest**



- c. Then I run my docker image to check my image formation occur correctly or not. I found my image is working ,For proof I have attached the following



screenshot:

```
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ docker run --rm -p 8080:8080 us-central1-docker.pkg.dev/my-sde-assignment/sde-docker-repo/sde-doc-image:latest
Unable to find image 'us-central1-docker.pkg.dev/my-sde-assignment/sde-docker-repo/sde-doc-image:latest' locally
latest: Pulling from my-sde-assignment/sde-docker-repo/sde-doc-image
6f172cdbcbef: Pull complete
4f8b24328d62: Pull complete
b17fff9d66ff6: Pull complete
486b373bcb47: Pull complete
f51a9e7b36b4: Pull complete
8deb7ac77d1b: Pull complete
57483073210f: Pull complete
Digest: sha256:45ef934d2d46b7d983b481a6efa4edbc3065668bce940177587fff25b8fa679e
Status: Downloaded newer image for us-central1-docker.pkg.dev/my-sde-assignment/sde-docker-repo/sde-doc-image:latest
Downloaded storage object images/IMG-20190901-WA0014.jpg from bucket my-sde-bucket to local file ./resources/IMG-20190901-WA0014.jpg
```

- Next I have created Google Kubernetes Engine name as **sde-assignment-cluster-1** and then connect to to GKE cluster.
- To create deployment, I have used the following command:  
**kubectl create deployment hello-app --image=us-central1-docker.pkg.dev/my-sde-assignment/sde-docker-repo/sde-doc-image:latest**
- Next I checked , my deployment is ready or not.

```
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ kubectl get deployment
NAME                    READY   STATUS    RESTARTS   AGE
hello-app-745fd7b56b-r4c7s 1/1     Running   0           24s
```

- On the time of deployment, It was in ready state but after some time it was not ready so I tried to do more deployment but it happens same with every deployment and it's reason **CrashLoopBackOff** as you can see in following screenshot :

```
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ kubectl get deployment
NAME                    READY   UP-TO-DATE   AVAILABLE   AGE
ai-app                  0/1     1             0           24h
api-facedetection       0/1     1             0           59s
detection-face          0/1     1             0           16h
face-detection          0/1     1             0           16h
hello-app               0/3     3             0           24h
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
ai-app-57794c5f9b-n6gfb             0/1     CrashLoopBackOff    17         24h
api-facedetection-787ff896f9-4kft9  0/1     CrashLoopBackOff    4          2m15s
detection-face-5cfbccc95d-zrjgk     0/1     CrashLoopBackOff    17         16h
face-detection-7995965f58-hlmtm     0/1     CrashLoopBackOff    17         16h
```

- Recently I have done one deployment name as api-face-detection and it also gone to unready state.

- To set the baseline number of deployment replica to 3 and to create HorizontalPodScaler resources , I have used the following command:

```
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 27h
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ kubectl scale deployment api-facedetection --replicas=3
deployment.apps/api-facedetection scaled
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ kubectl autoscale deployment api-facedetection --cpu-percent=50
horizontalpodautoscaler.autoscaling/api-facedetection autoscaled
kumar_281@cloudshell:~/python-vision/samples/snippets/face_detection (my-sde-assignment)$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
ai-app-57794c5f9b-n6qfb	0/1	CrashLoopBackOff	26	24h
api-facedetection-787ff896f9-4kft9	0/1	CrashLoopBackOff	14	48m
api-facedetection-787ff896f9-8h6xj	0/1	CrashLoopBackOff	5	3m24s
api-facedetection-787ff896f9-124rv	0/1	Completed	5	3m24s
detection-face-5cfbcccc95d-zrjgk	0/1	CrashLoopBackOff	26	17h
face-detection-7995965f58-hlmtm	0/1	CrashLoopBackOff	26	17h

- To generate kubernetes service for the api-face-detection used the following command:

```
kumar_281@cloudshell:~ (my-sde-assignment)$ kubectl expose deployment api-facedetection --name=api-facedetection-service --type=LoadBalancer
service/api-facedetection-service exposed
kumar_281@cloudshell:~ (my-sde-assignment)$ kubectl get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
ai-app-service	LoadBalancer	10.96.0.150	146.148.66.198	80:31246/TCP	25h
api-facedetection-service	LoadBalancer	10.96.3.154	<pending>	80:30707/TCP	17s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	27h

```
kumar_281@cloudshell:~ (my-sde-assignment)$ kubectl get service --WATCH
error: unknown flag: --WATCH
See 'kubectl get --help' for usage.
kumar_281@cloudshell:~ (my-sde-assignment)$ kubectl get service --watch
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
ai-app-service	LoadBalancer	10.96.0.150	146.148.66.198	80:31246/TCP	25h

- In the above screenshot , we got the external ip-address as **34.135.216.101** and we can see that its type is load balancer.

## REFERENCES

[https://cloud.google.com/vision/docs/face-tutorial?hl=en\\_US](https://cloud.google.com/vision/docs/face-tutorial?hl=en_US)

## YouTube Link

[https://www.youtube.com/watch?v=IRTUOGIO\\_T8](https://www.youtube.com/watch?v=IRTUOGIO_T8)



